


DOCKER concepts



docker

28/04/2017




Sommaire

1. Introduction
2. Installing docker
3. Introduction to images
4. Managing containers
5. Building images
6. Building images
7. Volumes
8. Networking
9. Docker machine
10. Docker compose
11. Docker Swarm
12. Security

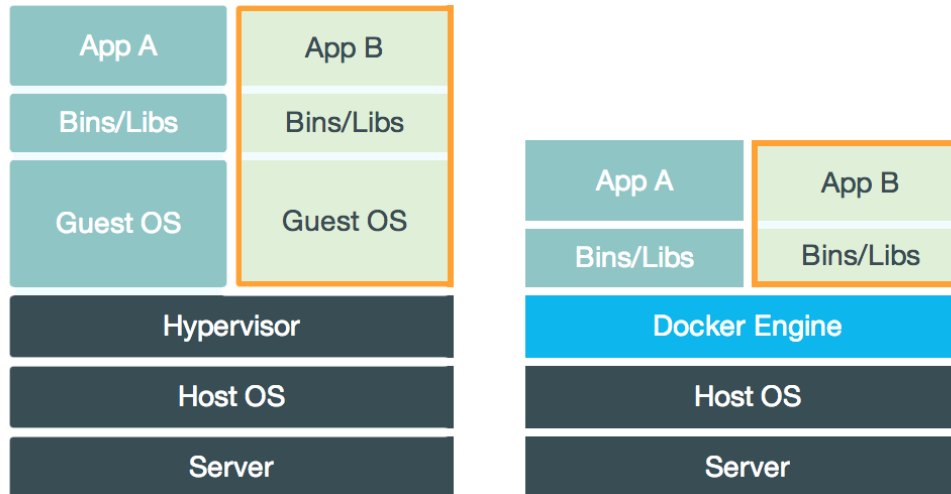
Introduction

Container based virtualization

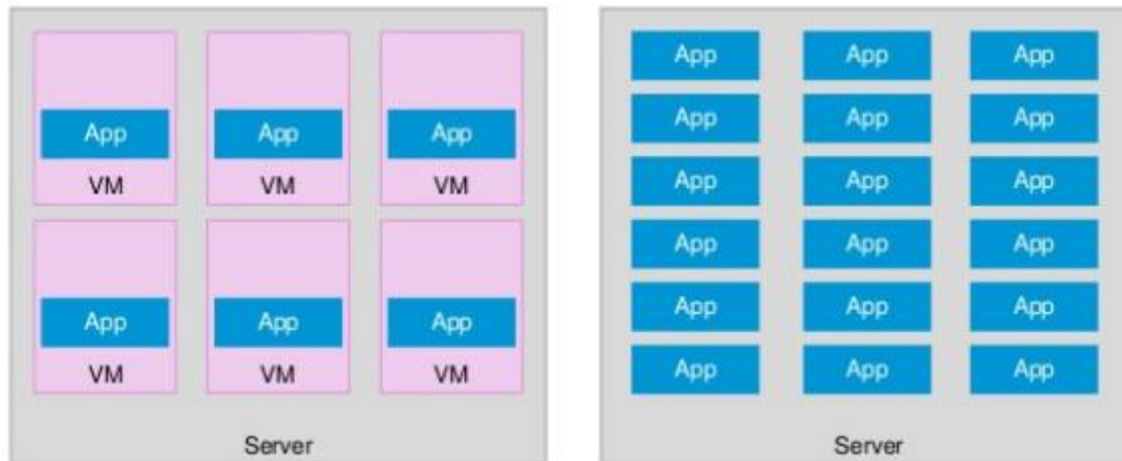
- ❑ Isolated systems
- ❑ Containers share the same OS kernel
- ❑ Containers hold the components necessary to run the desired software



Container based virtualization



Container based virtualization



Pets vs Cattle



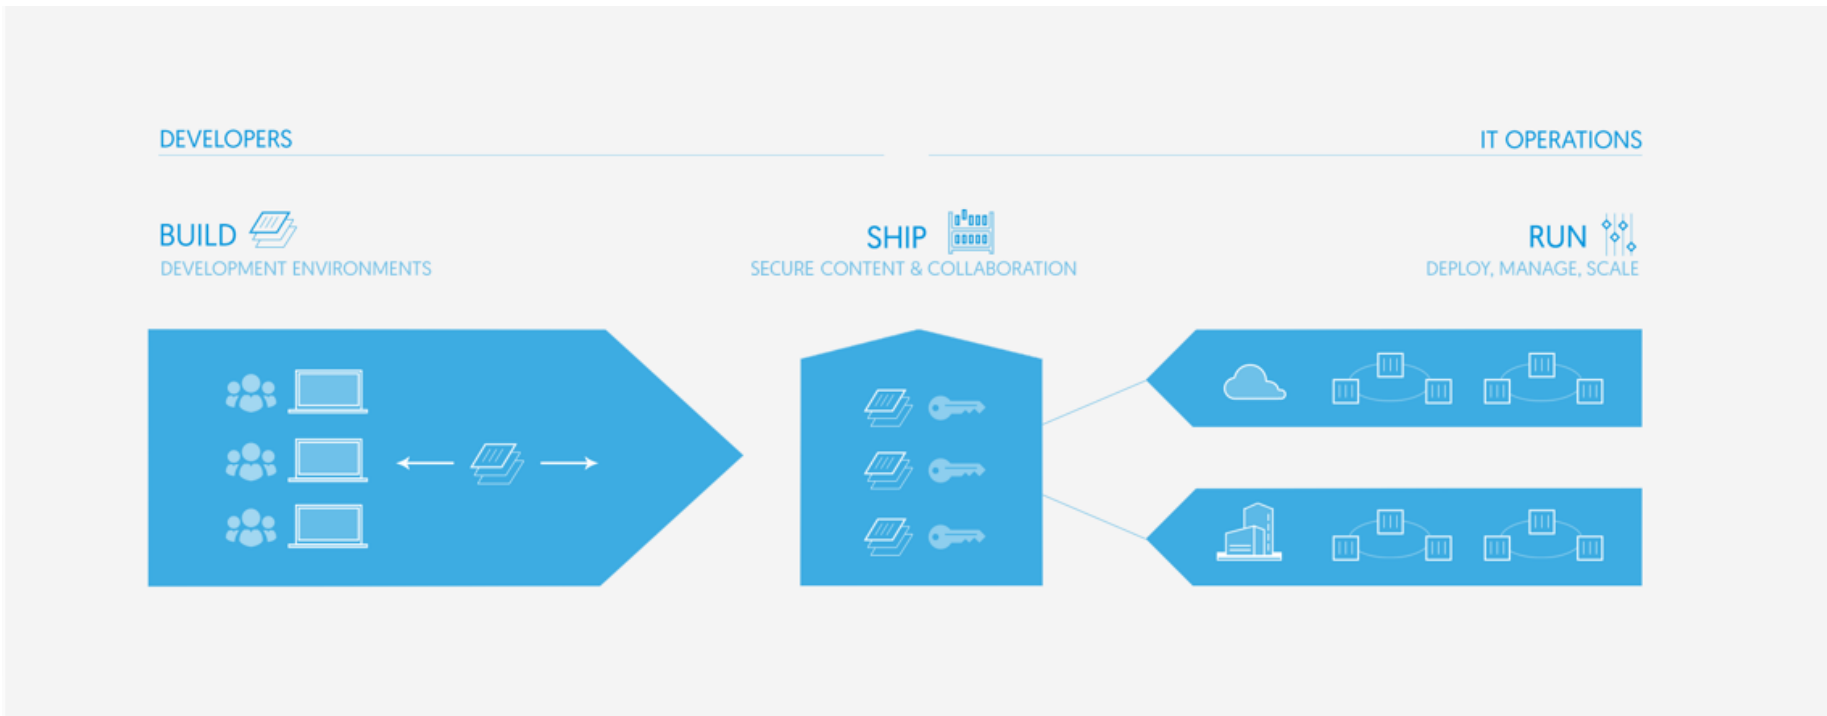
Pets vs Cattle



Docker Container

- ❑ Docker
 - ❑ API
- ❑ Docker container
 - ❑ cgroups and namespaces
 - ❑ Union file systems

Build Ship Run



Advantages of running docker

- ❑ Rapid application deployment (infra as code) + testing
- ❑ Portability across machines (hardware agnostic)
- ❑ Version control and component reuse (tag , layers)
- ❑ Sharing (registry)
- ❑ Lightweight footprint and minimal overhead
- ❑ Simplified maintenance (immutable infrastructure)

Drawbacks

- ❑ Be aware of security
 - ❑ Don't use it for monolithic application
 - ❑ You should respect some best practices
 - (<https://github.com/docker/labs/tree/master/12factor>)
-
- 1 - Codebase
 - 2 - Dependencies
 - 3 - Configuration
 - 4 - External services
 - 5 - Build / Release / Run
 - 6 - Processes
 - 7 - Port binding
 - 8 - Concurrency
 - 9 - Disposability
 - 10 - Dev / Prod parity
 - 11 - Logs
 - 12 - Admin processes

Installing Docker

Installing Docker

```
apt-get update
apt-get install -y apt-transport-https ca-certificates linux-image-extra-$(uname -r) linux-image-extra-virtual
apt-key adv --keyserver hkp://p80.pool.sks-keyservers.net:80 --recv-keys 58118E89F3A912897C070ADB76221572C52609D
echo "deb https://apt.dockerproject.org/repo ubuntu-xenial main" | sudo tee /etc/apt/sources.list.d/docker.list
apt-get update
apt-get install -y docker-engine
service docker start
usermod -aG docker ubuntu
```

Introduction to images

Differences bw images and containers :

- ❑ **Image**
 - inert, immutable, file
 - snapshot of a container
- ❑ **Container**
 - running instance of that image

Where to find docker images ?

- ❑ Docker Hub <https://hub.docker.com/>
 - ❑ Public Registry
 - ❑ Share images with co-workers
 - ❑ Manage images with tags
 - ❑ Use official Images



- ❑ Build your own and store it in a **private** registry



Managing containers


Intro to docker cli

```
docker --help
docker run hello-world
docker --version
docker pull
docker run
docker exec
docker inspect
docker rm
docker ps
docker stats
docker logs
...
```

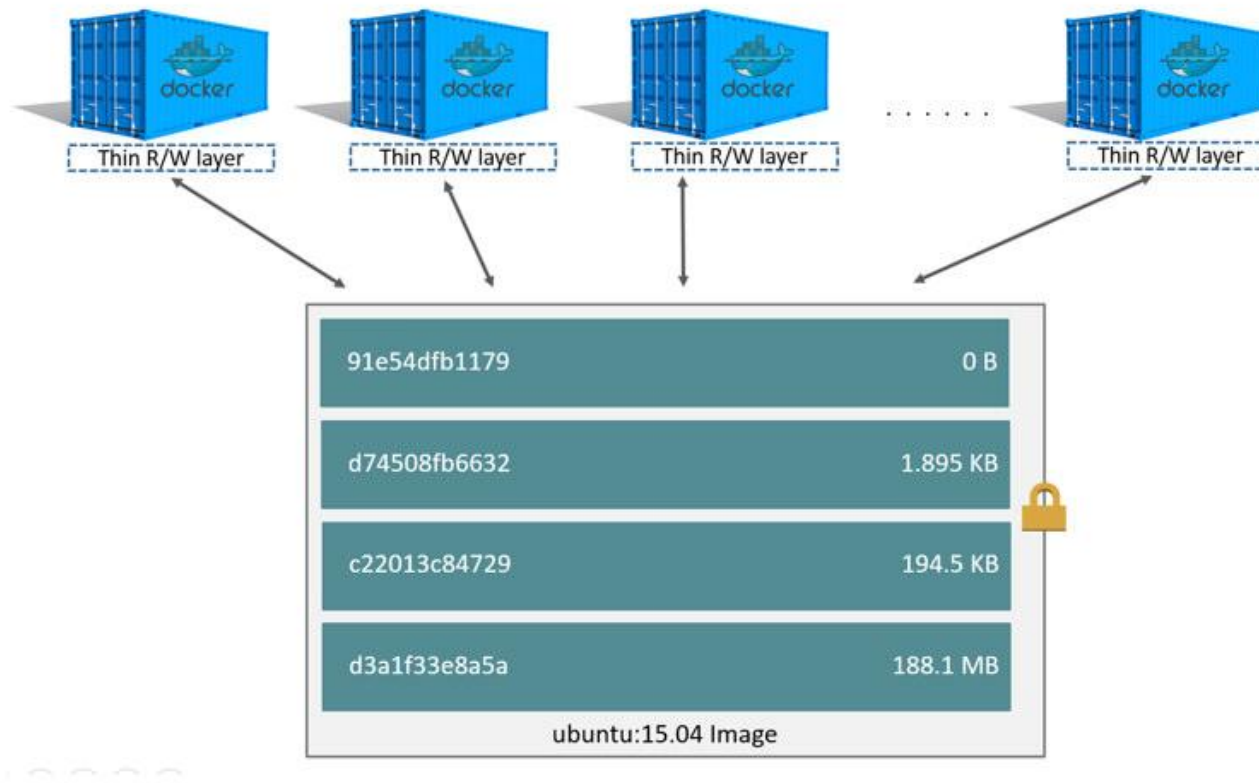


Building images

Images and layers



Images and layers



Images management

docker import / export

docker tag

docker push / pull

docker rmi

Dockerfile

- ❑ Docker can build images automatically by reading the instructions from a **Dockerfile**

```
docker build -t shykes/myapp .
```

- ❑ Instructions :

- ☐ FROM
- ☐ RUN
- ☐ ADD
- ☐ ENV
- ☐ EXPOSE
- ☐ LABEL
- ☐ USER
- ☐ WORKDIR
- ☐ VOLUME
- ☐ ...



Volumes

Data volumes

- ❑ Data persistence
- ❑ Share data between containers
- ❑ Plugins (host directory, shared storage, ...)

```
docker run -d -p 80:80 -v /home/ubuntu/localdir:/usr/share/nginx/html/ --name nginxsth nginxsth:latest
```



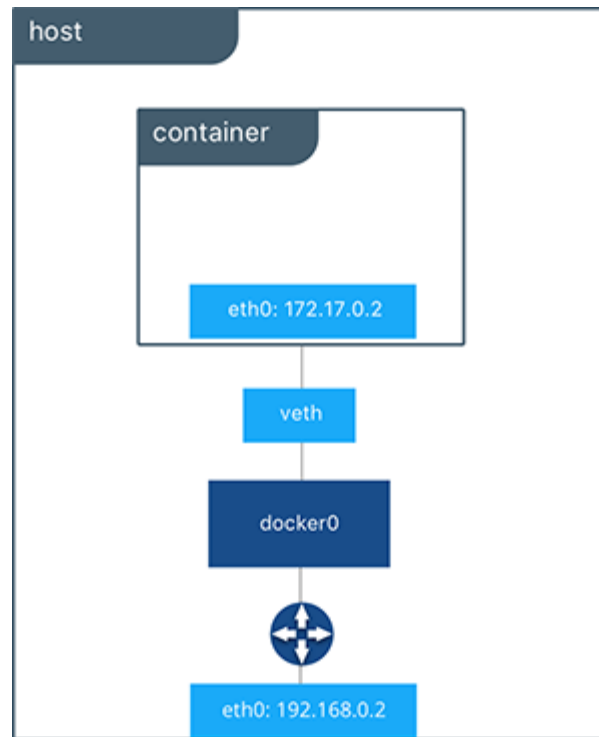
Networking

Port mapping

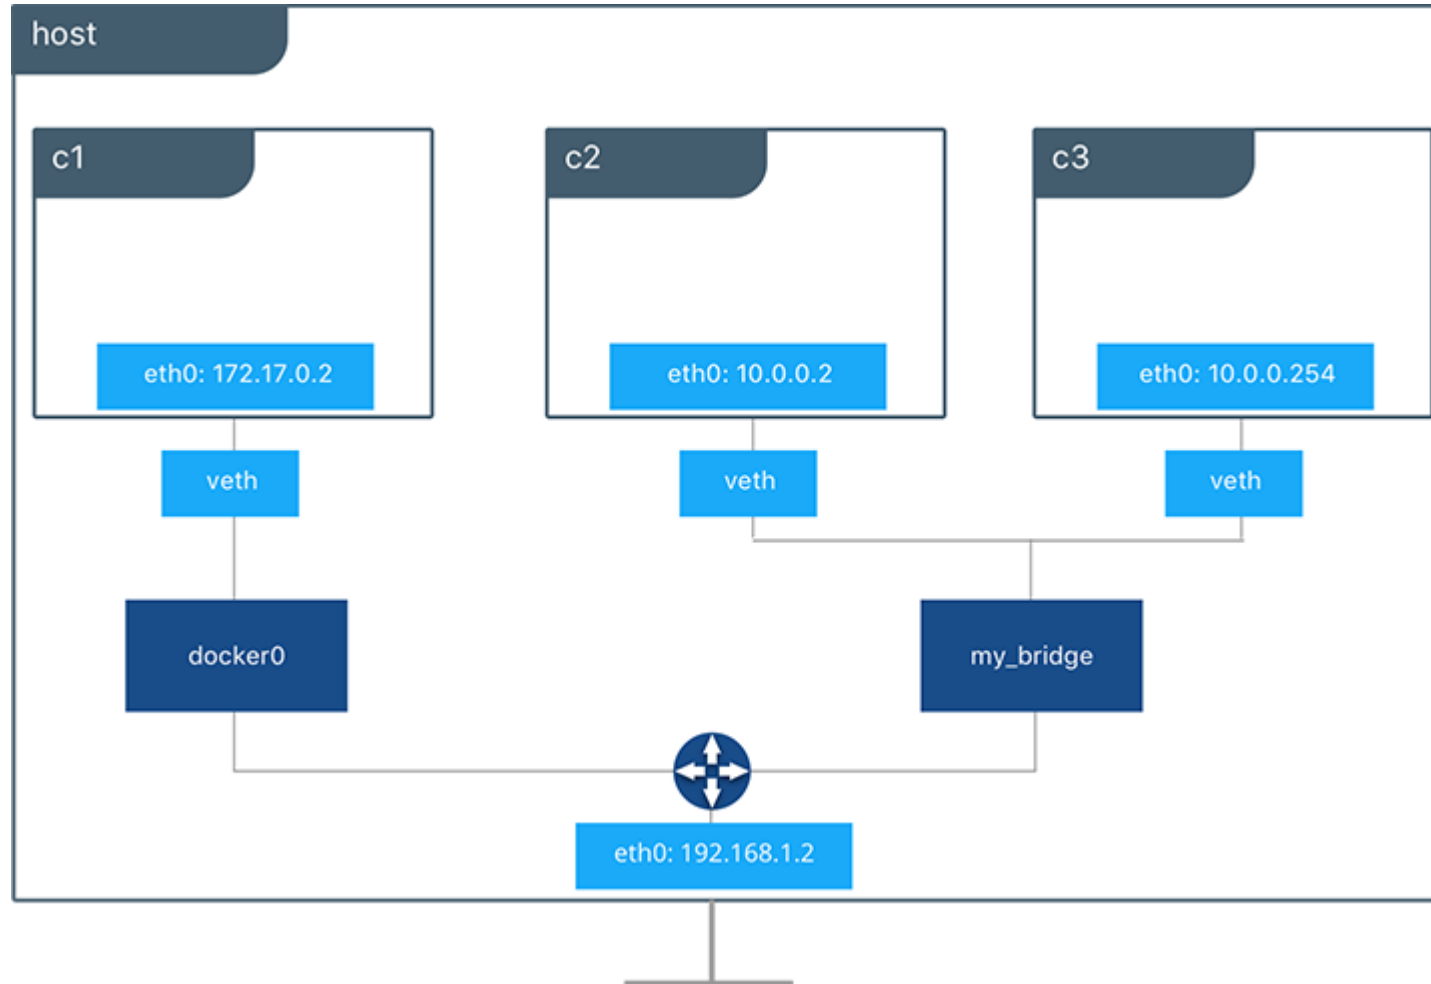
- ❑ Default:
 - ❑ Containers can make connections to the outside world
 - ❑ Outside world cannot connect to containers
- ❑ Mapping:
 - ❑ To accept incoming connections :
 - specify option -P or -p IP:host_port:container_port in 'run' command
 - ❑ iptables -t nat -L -n

- ❑ Embedded DNS
 - ❑ Embedded DNS to provide service discovery for containers (127.0.0.11:53)
 - ❑ Key/value store in Docker Engine
 - ❑ Network-scoped (Containers not on the same network cannot resolve each other's addresses)

Bridging



Bridging

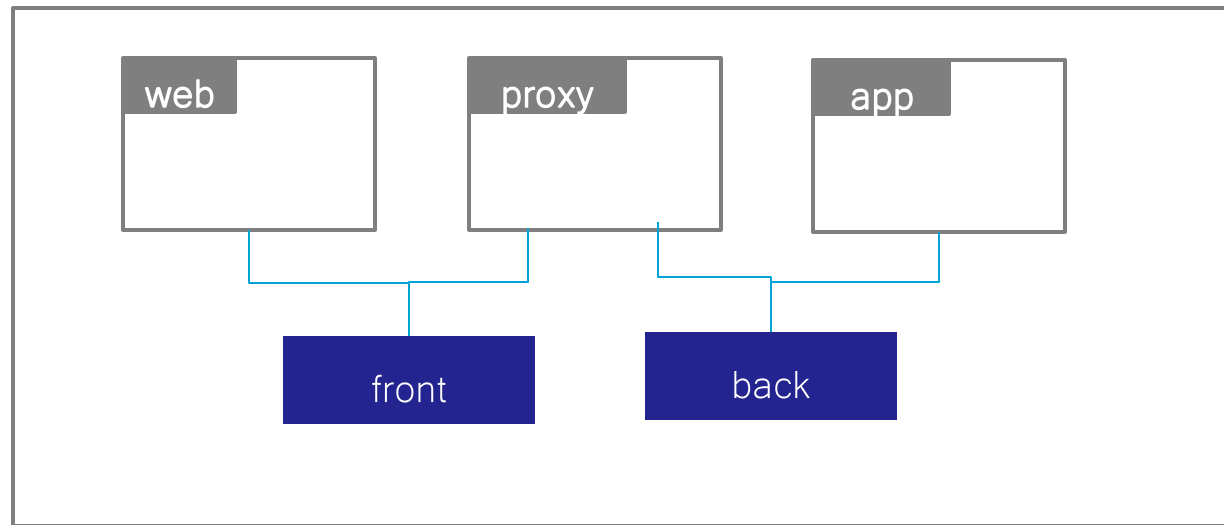


Network demo

```
docker network create --driver bridge isolated_nw
docker run -d -P --network=isolated_nw nginxsth:latest
docker network inspect isolated_nw
docker exec -ti fb82aad80f58 bash
ping 172.18.0.2 => ok
ping reverent_ramanujan => KO
docker run -d -P --network=isolated_nw --name c1 nginxsth:latest
docker run -d -P --network=isolated_nw --name c2 nginxsth:latest
docker exec -ti c1 bash
ping c2 => ok
```

Network demo


```
docker network create --driver bridge front
docker network create --driver bridge back
docker run -d --net=front --name web nginxsth
docker run -d --net=front --name proxy nginxsth
docker network connect back proxy
docker run -d --net=back --name app nginxsth
```




Multihost networking

- ❑ Overlay outside swarm mode
 - ❑ Requires a valid key-value store service
 - ❑ Consul, Etcd, and ZooKeeper
 - ❑ Configure docker engines to use the key-value store


- ❑ Overlay in swarm mode
 - ❑ Embedded key-value store




Overlay VXLAN



Overlay VXLAN



Overlay VXLAN






Docker machine

Docker machine

- ❑ Docker machine enables you to
 - ❑ Provision and manage multiple remote Docker hosts
 - ❑ Provision Swarm clusters
- ❑ Docker machine
 - ❑ Automatically creates hosts
 - ❑ Installs Docker Engine on them
 - ❑ Configures the docker clients (~/.docker)
- ❑ Drivers
 - ❑ AWS (ok)
 - ❑ Openstack (ok)
 - ❑ Virtualbox (ok)
 - ❑ Azure
 - ❑ Google Compute Engine
 - ❑ VMware



Docker compose

Docker compose

- ❑ Tool for defining and running multi-container Docker applications
=> compose file
- ❑ Create and start all the services
=> docker compose up

Docker compose

```
version: '2'

services:
  db:
    image: mysql:5.7
    volumes:
      - db_data:/var/lib/mysql
    restart: always
    environment:
      MYSQL_ROOT_PASSWORD: wordpress
      MYSQL_DATABASE: wordpress
      MYSQL_USER: wordpress
      MYSQL_PASSWORD: wordpress

  wordpress:
    depends_on:
      - db
    image: wordpress:latest
    ports:
      - "80:80"
    restart: always
    environment:
      WORDPRESS_DB_HOST: db:3306
      WORDPRESS_DB_PASSWORD: wordpress
volumes:
  db_data:
```

Docker Swarm

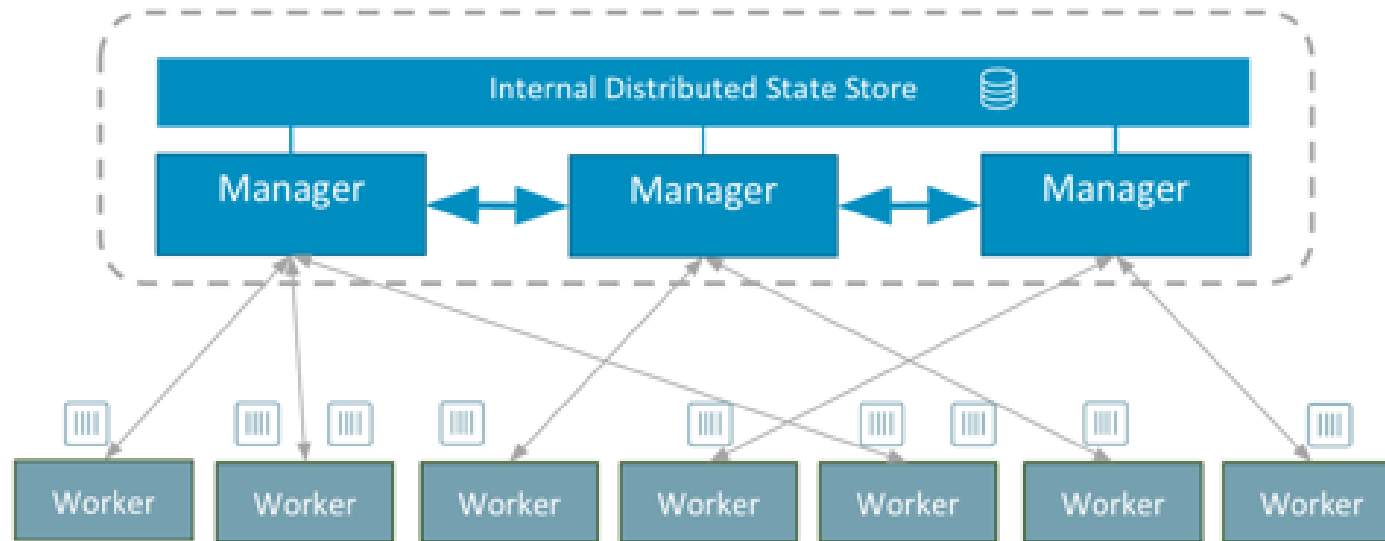
Docker Swarm

- ❑ Manage a cluster of Docker Engines:
 - ❑ Scaling
 - ❑ State reconciliation
 - ❑ Multi-host networking
 - ❑ Service discovery
 - ❑ Load balancing
 - ❑ Rolling updates
 - ❑ ...

Schedulling

- ❑ Scheduling Strategies
 - ❑ Spread
 - ❑ Binpack
 - ❑ Random
- ❑ Filters
 - ❑ Node constraint filter
 - ❑ Node name
 - ❑ Label
 - ❑ Storage driver
 - ❑ etc
- ❑ Container filters
 - ❑ Affinity
 - ❑ Label
 - ❑ etc

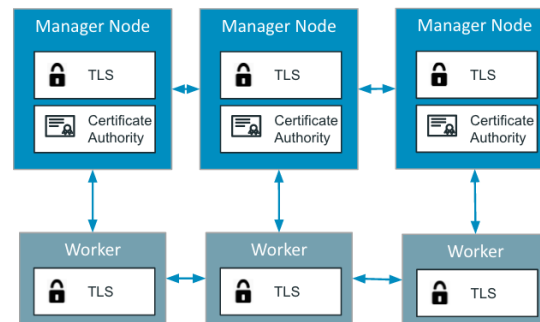
Architecture



Encryption

Control plane

- ❑ A CA is created with swarm init on the manager nodes
- ❑ All communication is encrypted over TLS.
- ❑ The node keys and certificates are automatically renewed (default 90 days)



Data plane

- ❑ IPSEC tunnels with AES algorithm between nodes. Keys are automatically rotated by managers every 12 hours.

```
docker network create --opt encrypted --driver overlay my-multi-host-network
```

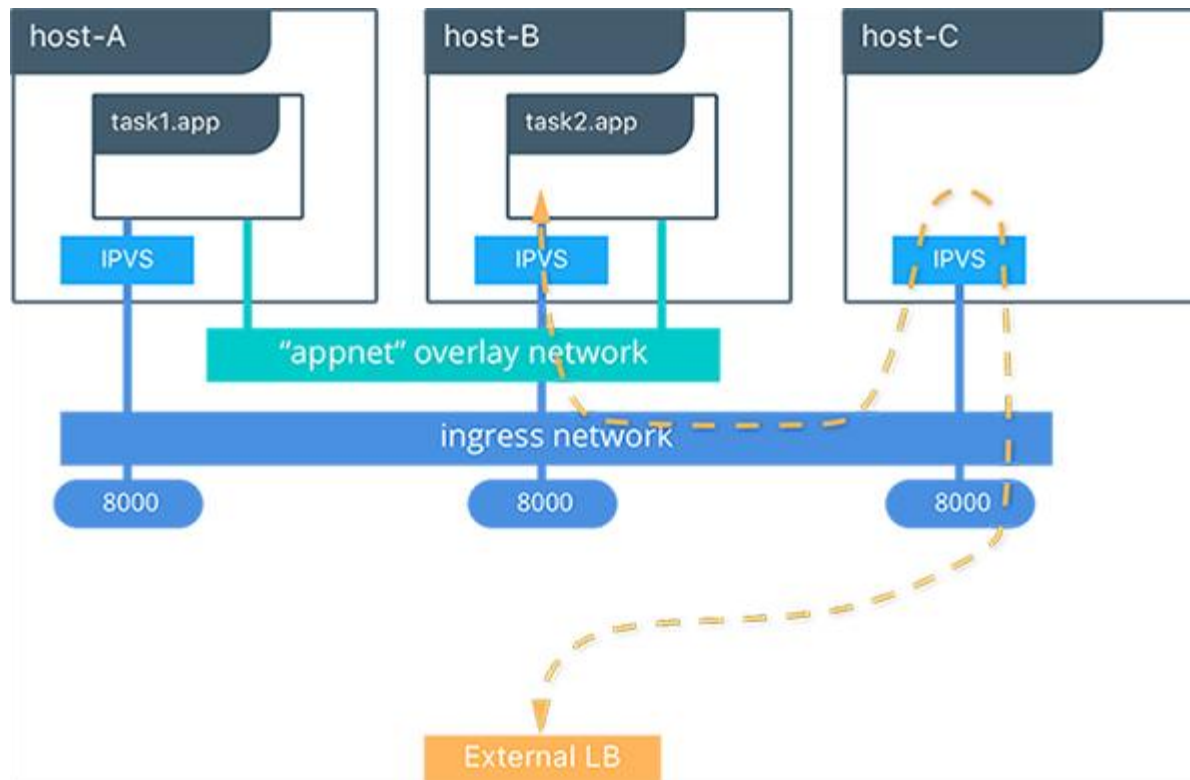
Internal Load balancing (from container to container)

- ❑ 2 options (--endpoint-mode)
 - ❑ DNS RR
 - ❑ IPVS Virtual IP (VIP)

```
docker service inspect myservice
[...]  
  
"VirtualIPs": [  
  {  
    "NetworkID": "a59umzkdj2r0ua7x8jxd84dhr",  
    "Addr": "10.0.0.3/24"  
  },  
]  
]
```

External Load balancing (to internet)

❑ Docker Routing Mesh



Stack

- ❑ Task = atomic unit ~ container
- ❑ Service = group of n tasks
- ❑ Stack = collection of services ~ application

```
services:
  lb:
    image: dockercloud/haproxy
    links:
      - web
    ports:
      - "80:80"
    roles:
      - global
  web:
    image: dockercloud/quickstart-python
    links:
      - redis
    target_num_containers: 4
  redis:
    image: redis
```

Secret management

- ◆ Use it for data that should not be
 - transmitted over a network
 - stored unencrypted in a Dockerfile or in your application's source code.
- ◆ such as a
 - Password
 - SSH private key
 - SSL certificate
 - etc

```
docker secret create
docker secret inspect
docker secret ls
docker secret rm
--secret flag for docker service create
```



Security

Security

- ❑ SANS Institute checklist:
 - ❑ Ensure good host security
 - ❑ Check Image Provenance
 - ❑ Monitor Containers
 - ❑ Do Not Run Container Processes as Root
 - ❑ Do Not Store Secrets in Containers
 - ❑ Base Image Security
 - ❑ Limit container resources

<https://www.sans.org/reading-room/whitepapers/auditing/checklist-audit-docker-containers-37437>



News

New in 1.13

- ❑ Compose to deploy stack in Swarm
- ❑ New cli commands : docker container and docker image
- ❑ Docker system command
- ❑ Secret management

- ❑ Experimental :
 - ❑ Docker service logs
 - ❑ Docker Metrics in Prometheus format